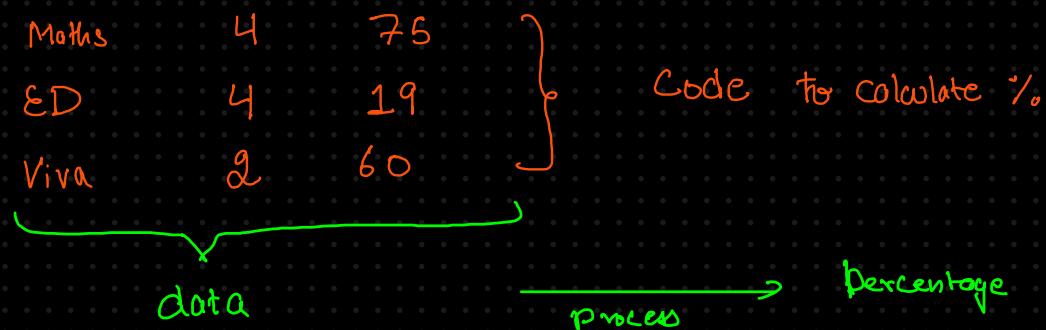


Map Reduce

NSIT → marks ⇒ percentage



Mayork
II



1



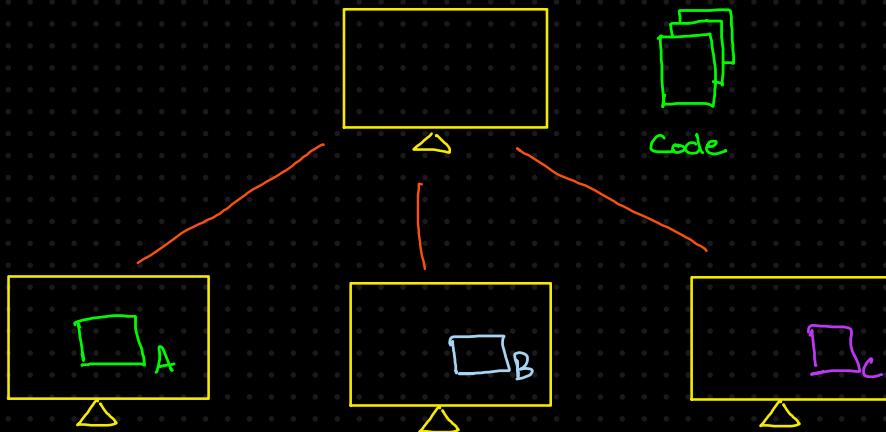
F2



F3

Google chrome
Voting

Delhi



data → code

* data locality
+ parallel processing

~ Divide & Conquer
Parse digm.

Map Reduce

GMR \longrightarrow MR

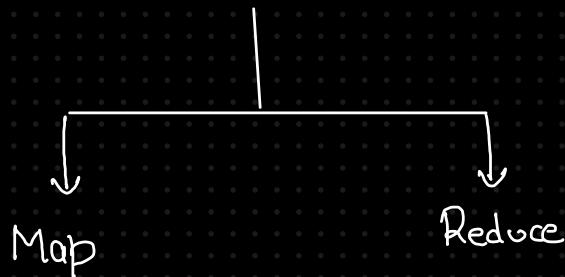
1. it is difficult to write
2. Theory is a lot imp \longrightarrow distributed processing
Practical is not that imp.

Map Reduce is a programming model designed for processing large dataset in a distributed computing manner

- Scalability
- Fault tolerance
- parallelism

Smaller task

1. We will have data distributed Parallelly via hdfs which is required

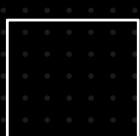
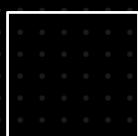


input data is split into chunks, & each chunk processed by a mapper

\downarrow
Mapper produces intermediate output



Reducer combines it



A

B

C

D

5 min

5 min

5 min

5 min

= 20 mins



5 min

5 min

5 min

5 min

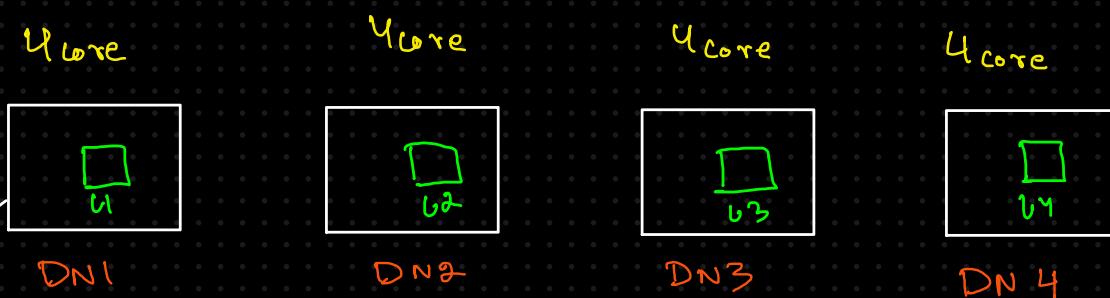
= 7 min



2 min

Each and every Input or Output of MR job
is a Key-value pair.

Say we have a 400 mb file in a hadoop cluster



1 core can process 1 mapper job on 1 block/input
of data

4 cores

it can parallelly process 4 blocks/IS of data

total 16 cores

1 file → 32 blocks
128mb block size

so in this case as blocks >> cores, we will process
the job in batches.

Word Count example

1. Input → text with sentences

2. Mapper → Value (`'hello'`, `'world'`)
↓

(`('hello', 1)` (`'world'`, 1))

3. Combiner (optional)

4. Shuffle and sort

5. Reducers

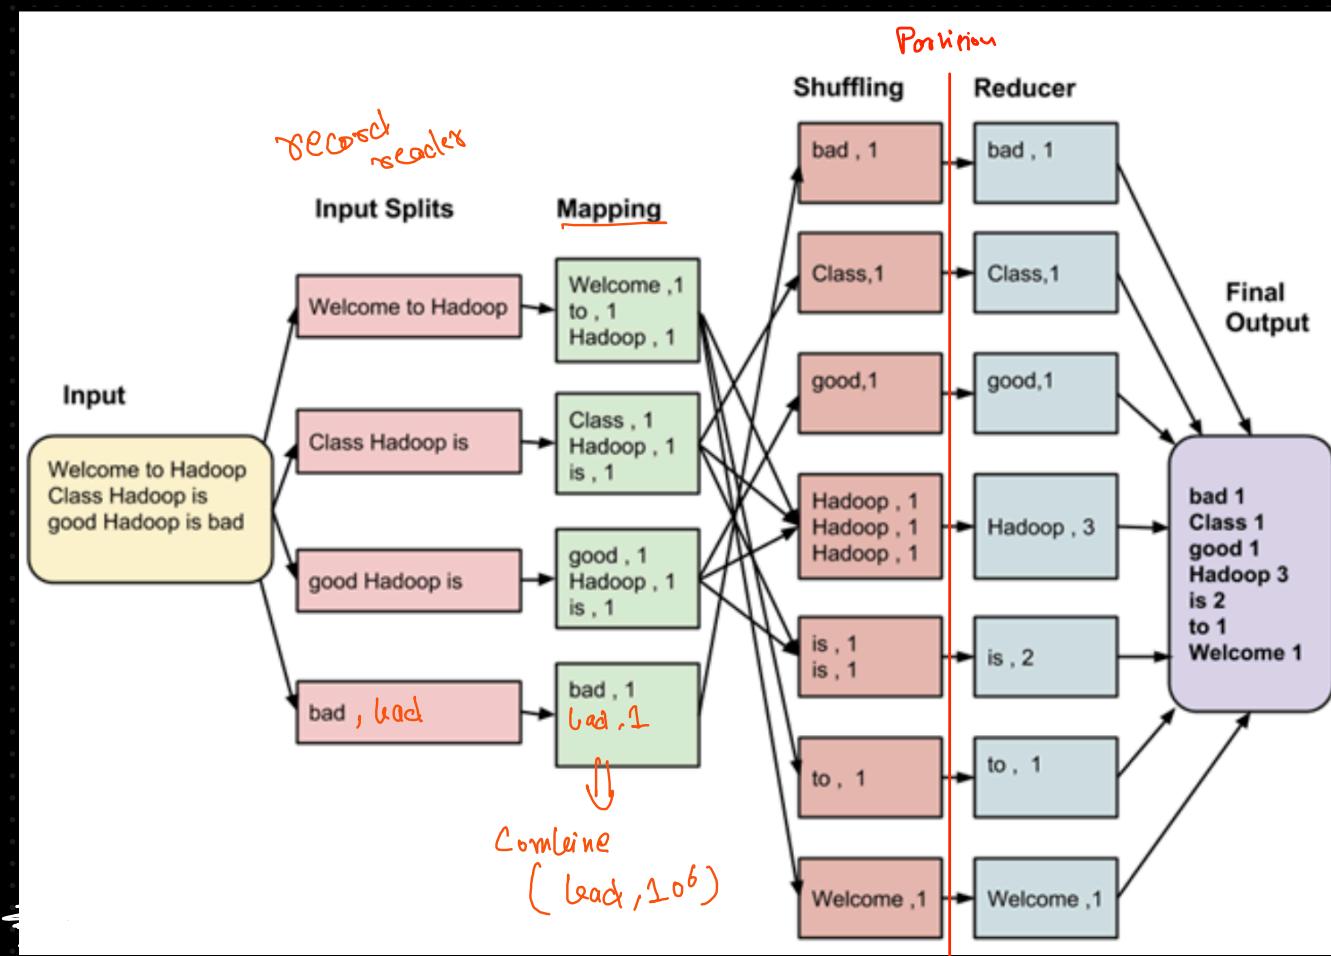
Aggregate the values

Goku, {1, 1, 1, 1, 1}



(Goku, {43})

6. Final Output



Gohu vegeta Gohou → (Gohu,1) (vegeta,1) (Gohou,1)

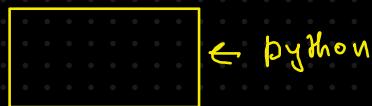
Gohou Frieza Gohu Gohu → (Gohou,1) (Frieza,1) (Gohou,1)
(Gohu,1)

↓ Combiner

(Gohu,2)

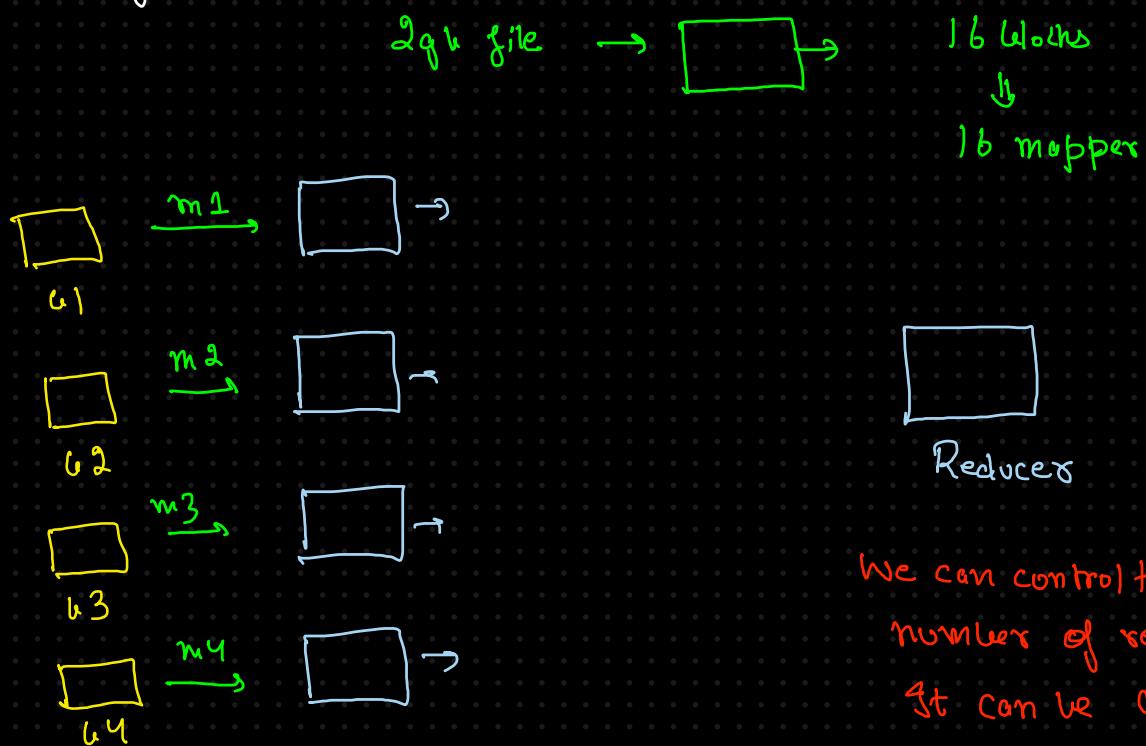
Step	Process	Output Example
Input Splitting	Splitting input file into HDFS chunks.	Split 1: Goku Vegeta Gohan , Split 2: Frieza Goku , etc.
Mapper Phase	Generating key-value pairs (word , 1) from each split.	("Goku" , 1) , ("Vegeta" , 1) , ("Gohan" , 1)
Combiner Phase	Optional local aggregation at Mapper to reduce shuffle data.	("Goku" , 2) , ("Frieza" , 1)
Shuffle and Sort	Grouping and sorting intermediate key-value pairs by key.	("Goku" , [1 , 2 , 1 , 2]) , ("Frieza" , [1 , 1 , 1])
Reducer Phase	Aggregating grouped data to produce final results.	("Goku" , 6) , ("Frieza" , 3)
Final Output	Writing final aggregated results to HDFS.	Frieza 3 , Gohan 3 , Goku 6 , Vegeta 2

$(\text{word}, 1)$
 $(\text{word}, 1) \Rightarrow 10^6$ ↗
 \vdots
 $(\text{word}, 1)$



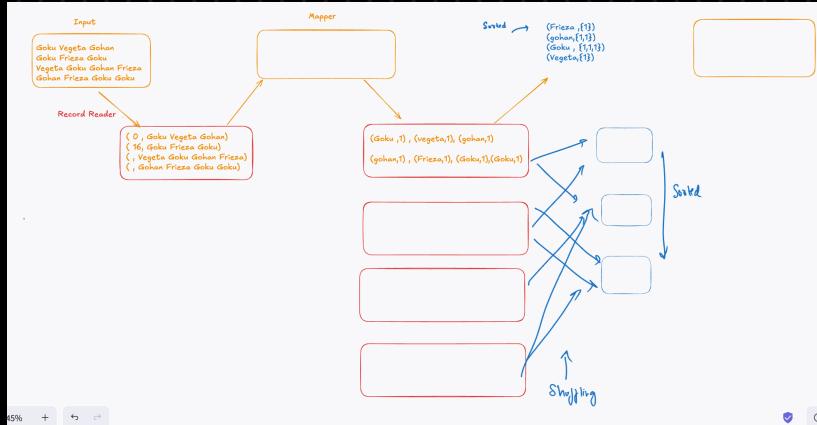
MR doesn't make sense for smaller data.

No. of Mappers ? = number of blocks



We can control the
number of reducers
It can be 0, 1, ... n





12th Jan

Map Reduce → 2 Reducers

0 reducers

log file → 2 blocks

Input split

YARN

Spark



PySpark

2. Map Reduce with more than 1 reducer

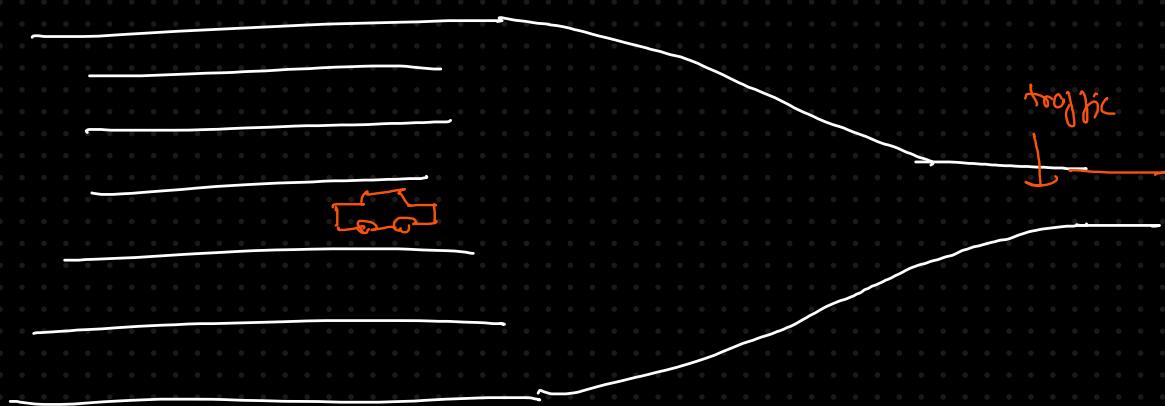
flipkart, amazon

iphone → google

{ flipkart → iphone, fm
amazon → iphone, pay,
myntra → cosmetics,

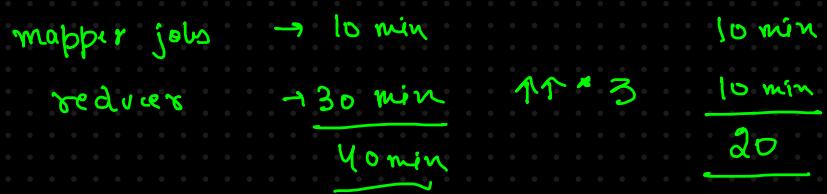
} billions of these rows

iphone → [amazon
flipkart
apple.com
...]



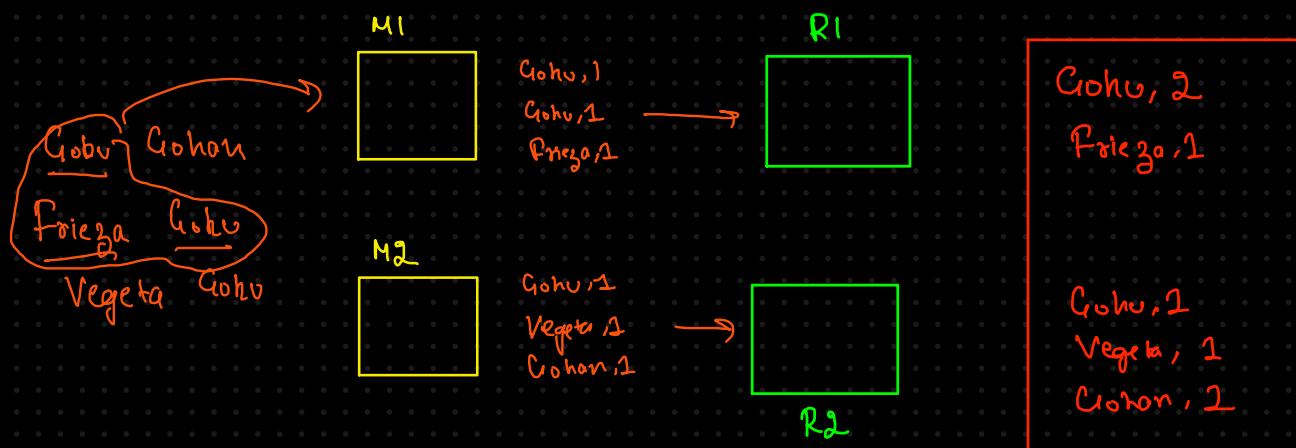
~~Always~~ Always make sure that mappers handles the heavy lifting in terms of processing

Let us have more than one reducers



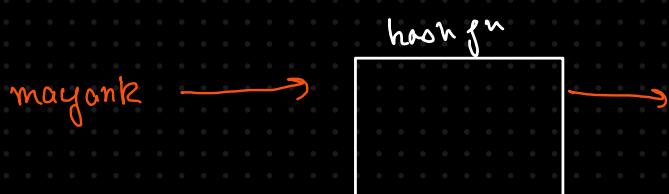
Partitioning

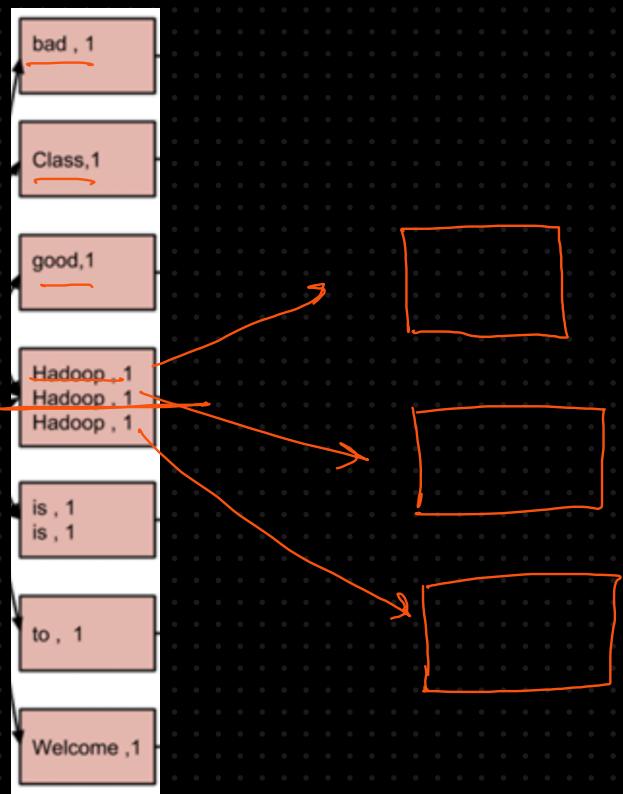
We can not have chain of reducers



Same keys should always go to same reducers

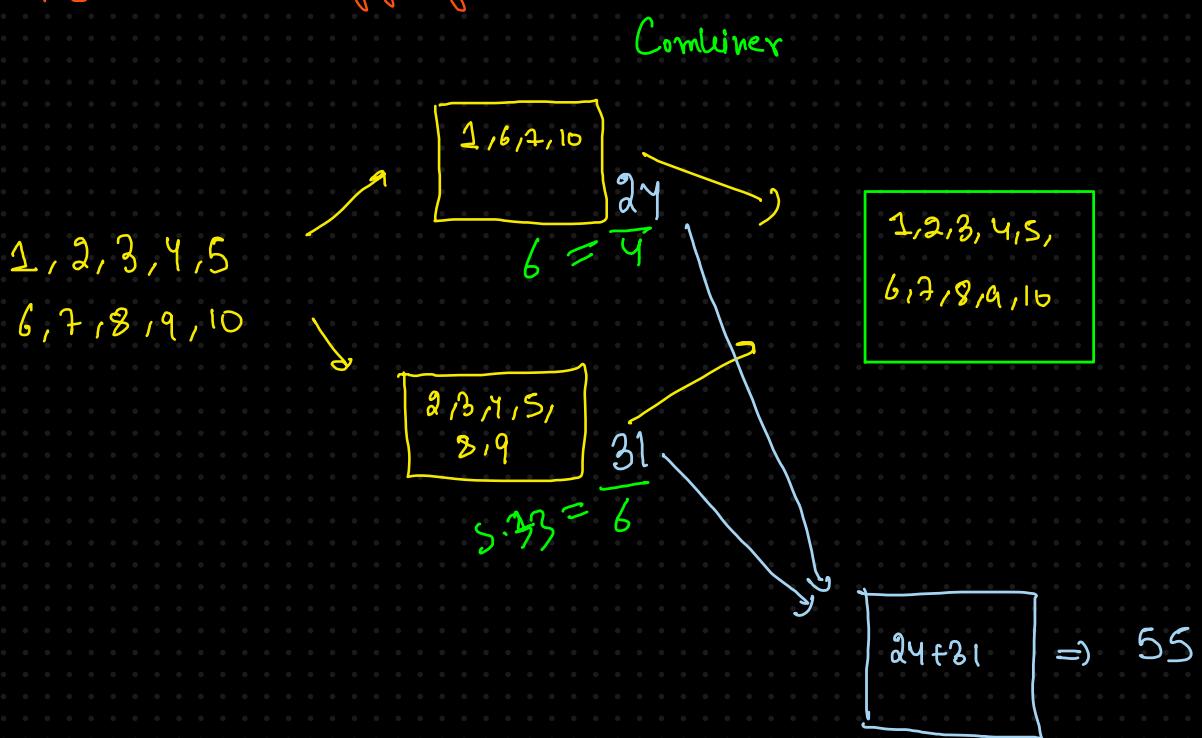
% modulus a very good hash function





Combiner

We do local aggregation



Then, can we always use combiner?

Non associative or non commutative approach

e.g. org, median

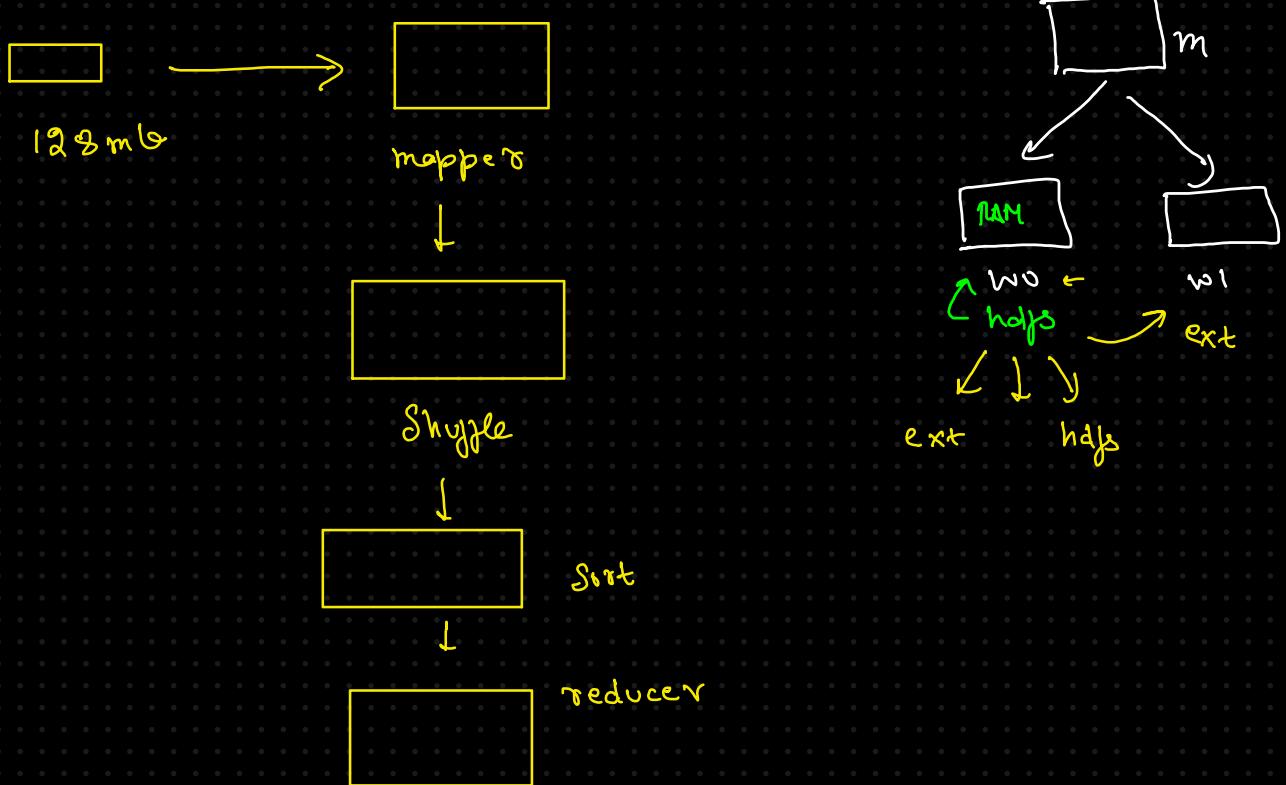
Map Reduce with zero Reducers

gohu 1

Examples → format conversion

Mapper O/P is final output

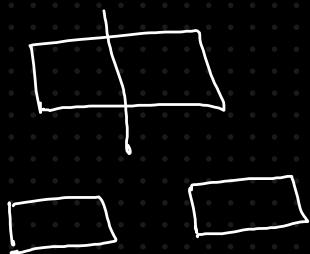
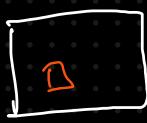
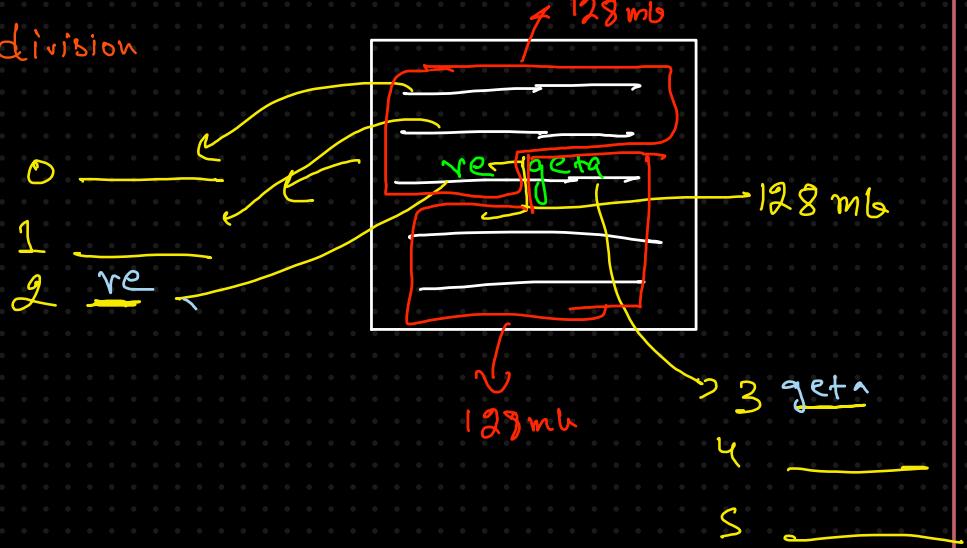
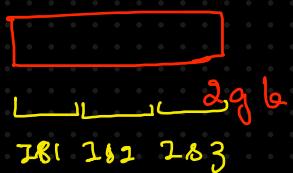
Map Reduce



Input splits

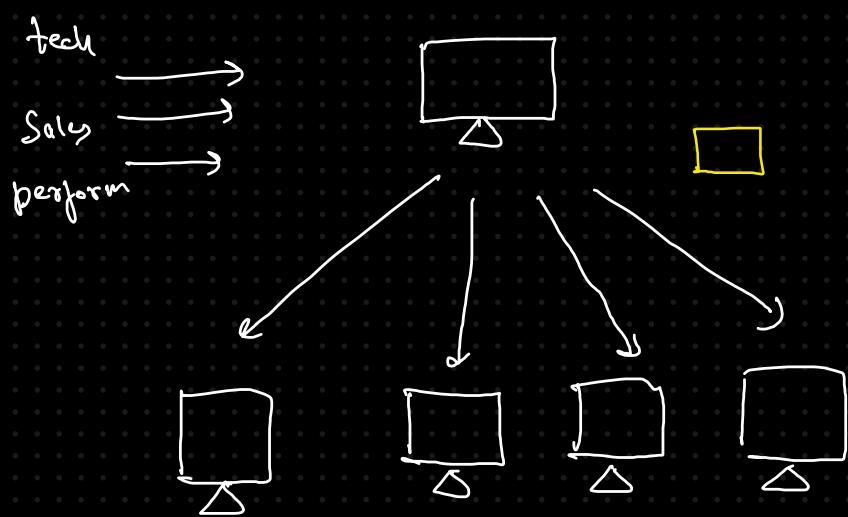
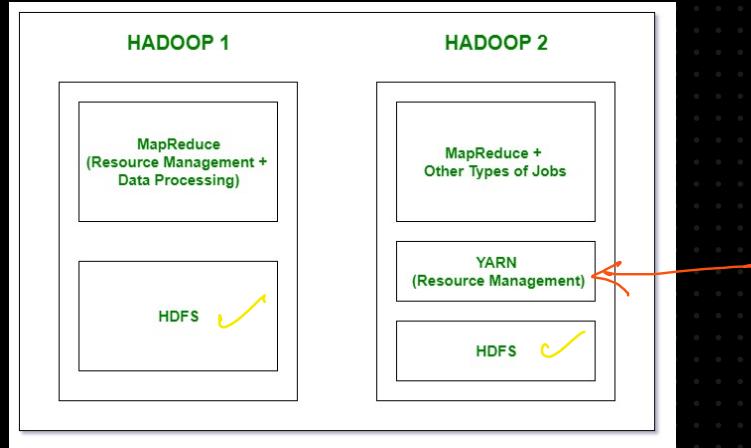
blocks → physical division

input split → logical division



Record

YARN (Yet another resource Negotiator)



JP1 → Name node
JP2 → Data node
JP3 → SNN
JP4 → RM
JP5 → NM

Client 1



Client 2

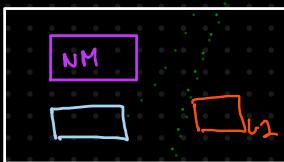


Client 3



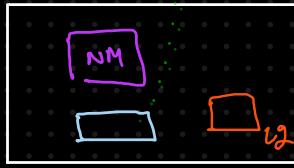
DN1

Word Count



DN2

App master



DN3

RM jobs

1. Scheduler
2. AM (application manager)

Container

Step by step

1. Submission of Application → RM accepts and creates a App Manager
↓
App master
2. request of resources
3. Allocation of resources → Scheduler
4. Launching of Container → RM communicate with NM
5. Task execution
5a. Monitoring & Fault Tolerance
6. Completion → App master will free up the resources
App master also killed & release its resources
7. Success →

